

Exercise 3: Running parallel image processing workload with Azure Batch

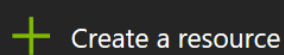
Duration: 45 minutes

In this exercise we will execute a simple ImageMagick resizing job on a set of input jpg files processed in parallel using Azure Batch service.


We will use Batch Labs which is an open-source GUI application for managing Azure Batch accounts, designed to make the use of Azure Batch service easy and straight-forward.

Step 1: Provision Azure Batch account

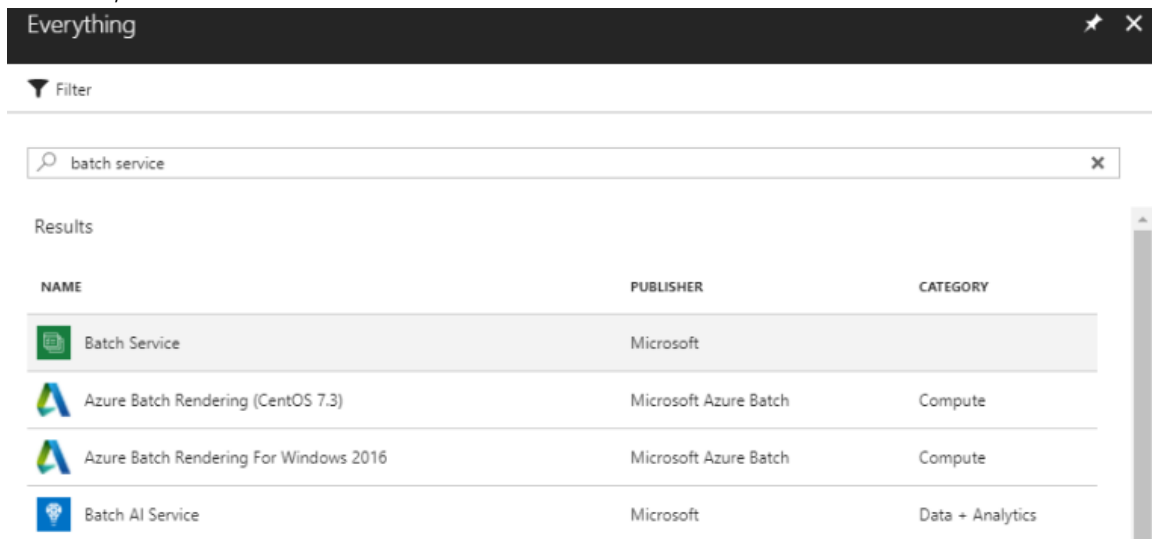
1. Navigate to Azure Portal in the browser.
2. Select **+ Create a resource**.



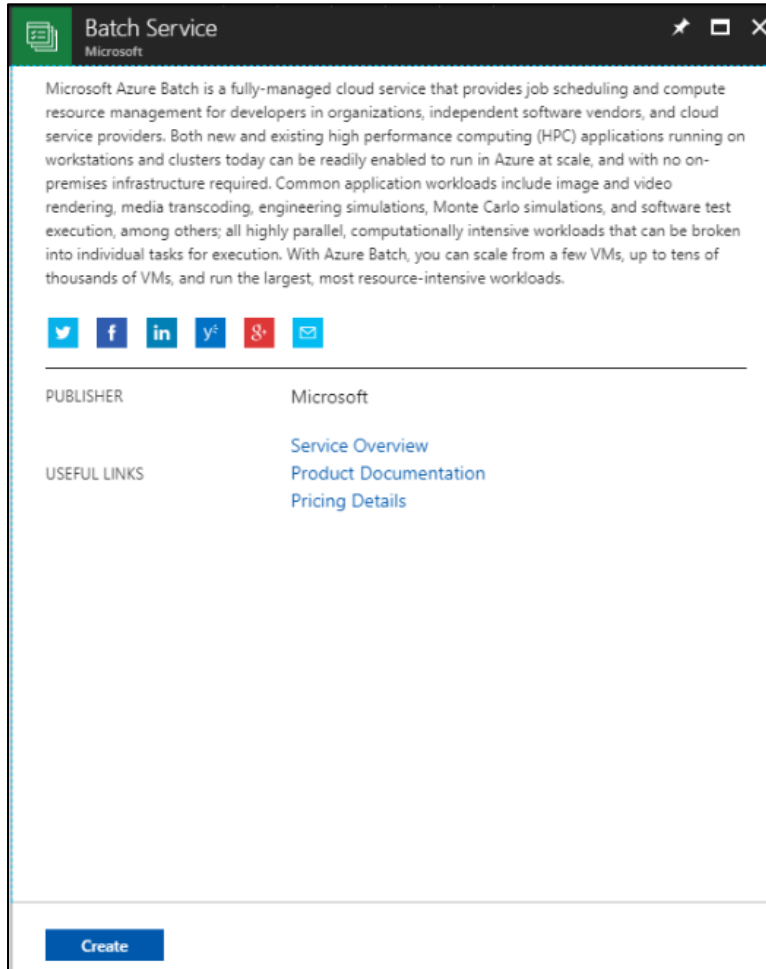
3. Search the Marketplace for **Batch Service**.



4. In the list, select **Batch Service**.

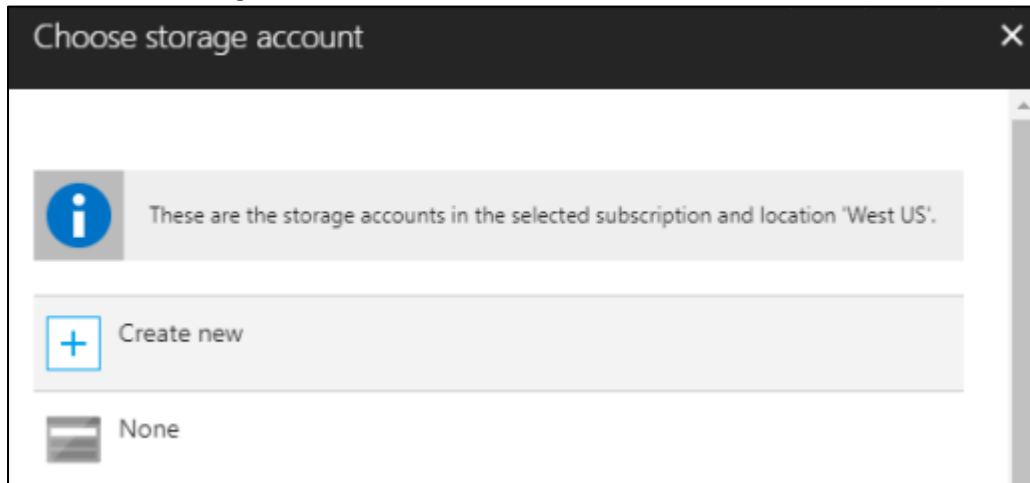


5. Select **Create** on the Batch Service blade.



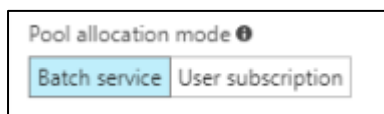
6. On the New Batch Account blade, specify the following:
- Account name:** Provide a name for your new Batch Account. The name you choose must be unique within the Azure region where the account is created (see Location below).
 - Subscription:** Select the subscription in which to create the Batch account.
 - Resource group:** Select a new resource group and name it (e.g. batchdemo).
 - Location:** The Azure region in which to create the Batch account.

7. Select **Storage account**.
8. In the Choose storage account, select **Create new**.



9. On the Create storage account blade:
 - a. **Name:** Provide a unique name for the new Azure Storage account that will be attached to your Batch Service. E.g. use scfelab<n> where n is the number on your desk.
 - b. **Performance:** Leave this at Standard.
 - c. **Replication:** Leave this at Locally-redundant storage (LRS).
 - d. Select **OK** at the bottom of the blade.

10. Back on the New Batch Account blade, leave the Pool allocation mode set to **Batch service**.

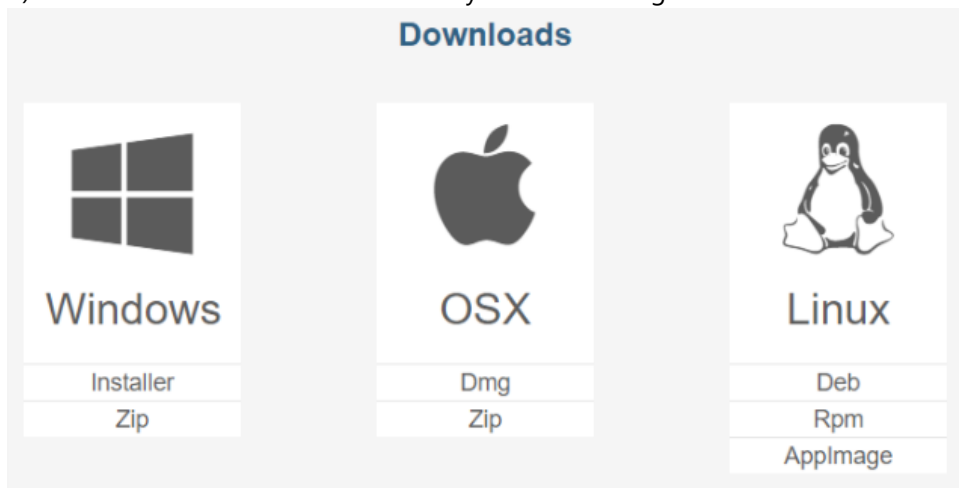


11. Select **Create** to create the new Batch account. The provisioning should take around 1 minute.

Preparation step 2: Install BatchLabs

1. Using a web browser on your local machine, navigate to <https://azure.github.io/BatchLabs/>.
2. Scroll down until you see the **Downloads** section.

3. Select the download appropriate to your OS (e.g., select Installer for Window or DMG for OS X). The instructions that follow assume you are installing on Windows.

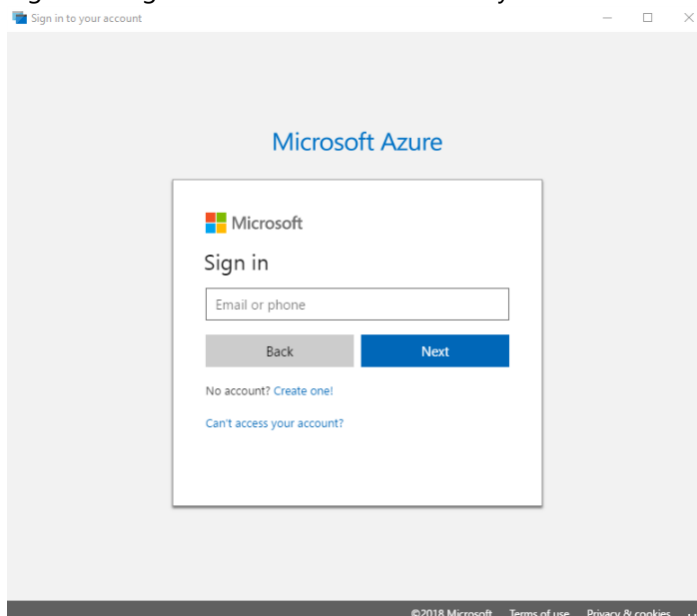


Further instructions assume that **Batch Labs 0.13.1 (beta)** is installed.

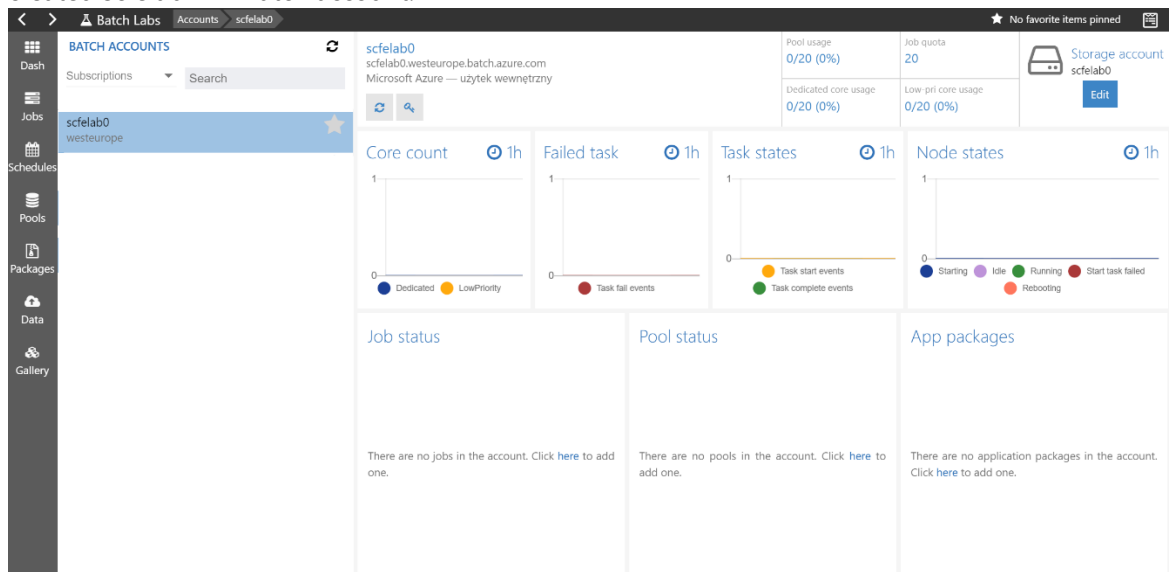
4. Run the installer.



5. Sign in using the account associated with your Azure Subscription.



6. Wait for Batch Labs application to complete loading. When the main Batch Labs window opens, in the left panel you should be able to select your Azure subscription and the just created scfelab<n> Batch account.

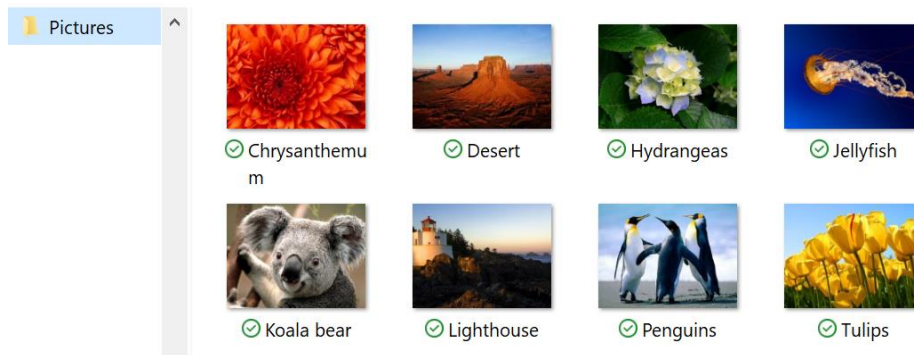


7. Your installation of Batch Labs is ready for use.

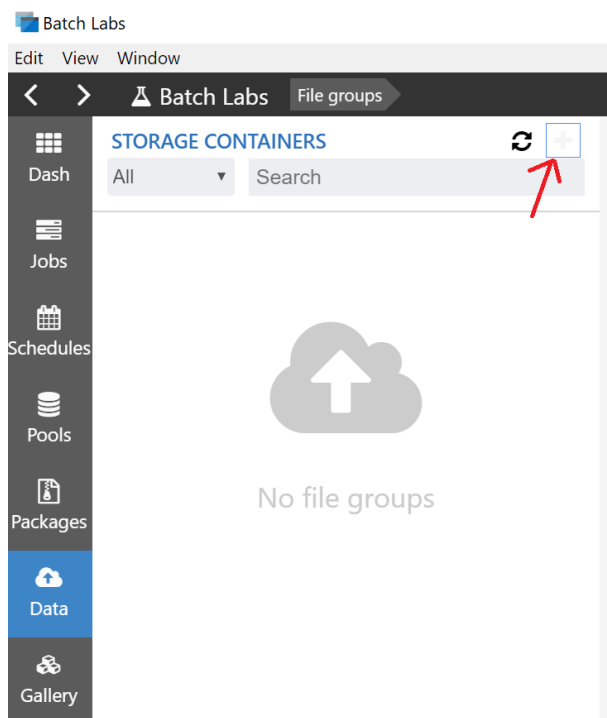
Now we will use the just installed Batch Labs application to process in parallel a set of input jpeg files with ImageMagick application.

Step 1: Upload input files to Batch Labs input file group

1. Download sample set of input jpeg files from <https://github.com/tojozefi/scfelab/blob/master/Pictures.zip>
2. Uncompress the content to any local folder, preserving the internal directory structure – folder 'Pictures' with 8 jpeg files.



3. Navigate to **Data** panel in Batch Labs and click **Add file group (+)** to create a new storage container:



4. Select **From local folder (file group)** from the drop-down list:

Empty container

Empty file group

From local folder(File group)

5. Fill out the **Create file group** form that opens:

Create file group

Upload files into a managed storage container that you can use for resource files in your jobs and tasks

1 General info

Basic information about the file group

File group name

image-magick-input

2 Files

Choose a directory containing the files you want to include in the file group

Select directory

Selected folder details

C:\Users\tojozefi\Desktop\Pictures



Include sub directories

3 File options

How data should be laid out in storage

Prefix

Flatten



Upload

Upload and close

Close

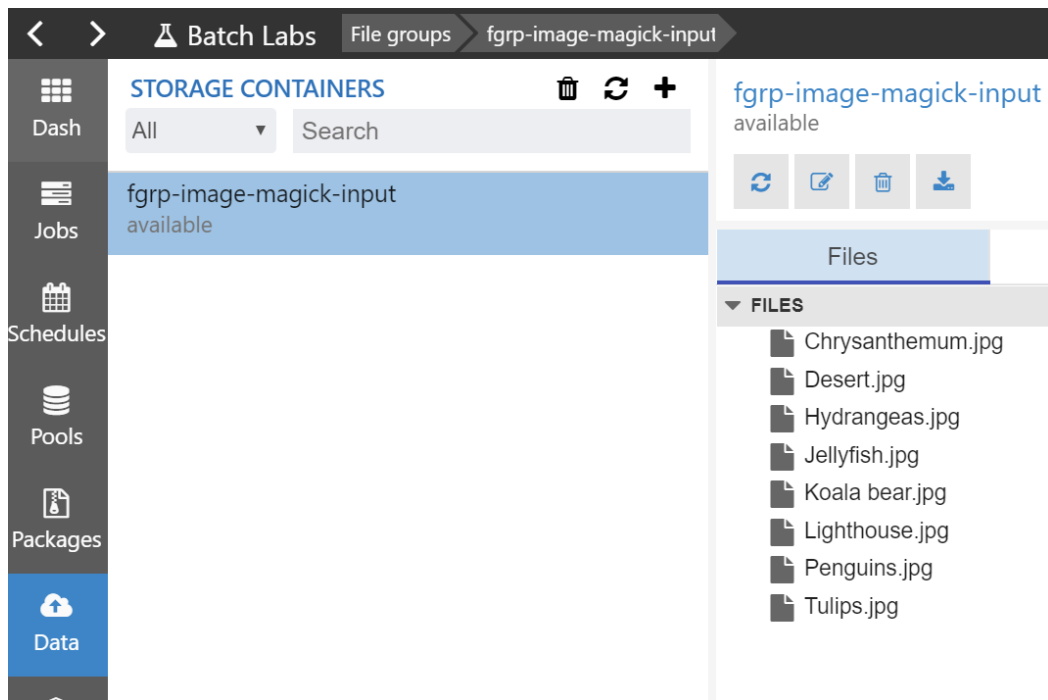
In (1) **General info** provide the name for the input file group, e.g. **image-magick-input**

In (2) **Files** select the *Pictures* directory with sample jpeg files from local disk and check **Include sub directories** box.

Leave all options in (3) **File options** intact.

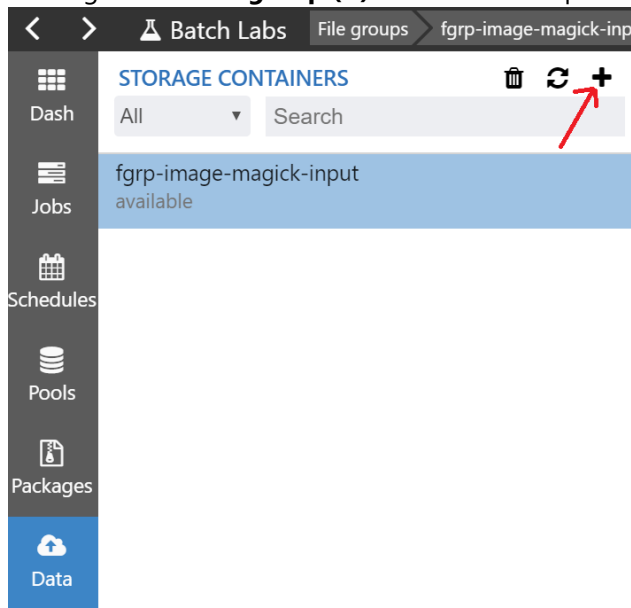
6. Click **Upload and close** button and wait for the file upload to complete.

On successful upload you should see all jpeg files in the **fgrp-image-magick-input** file group in Batch Labs Data panel:



Step 2: Create output file group

- Click again **Add file group (+)** to create an output storage container:



- Select **Empty file group** from the drop-down list:

Empty container

Empty file group

From local folder(File group)

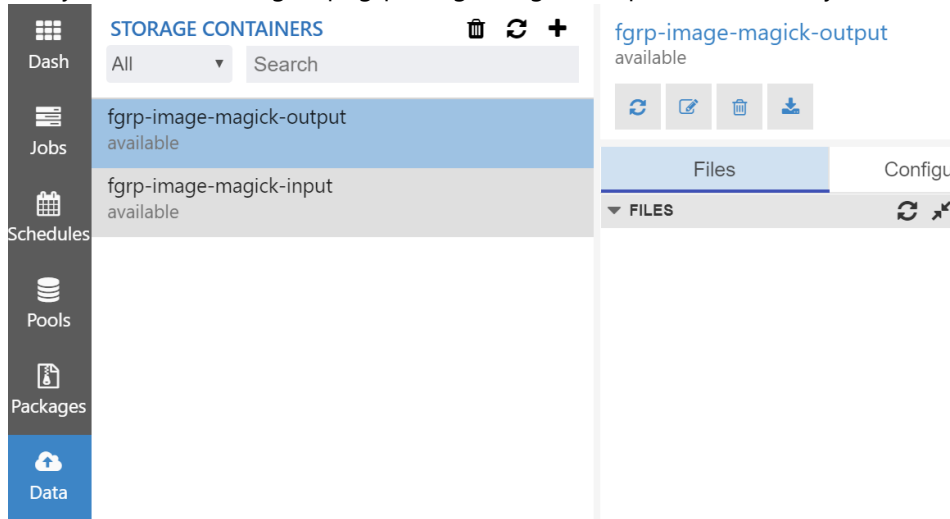
9. Provide a name for the output file group, e.g. **image-magick-output** and confirm.

Create a new empty file group

image-magick-output

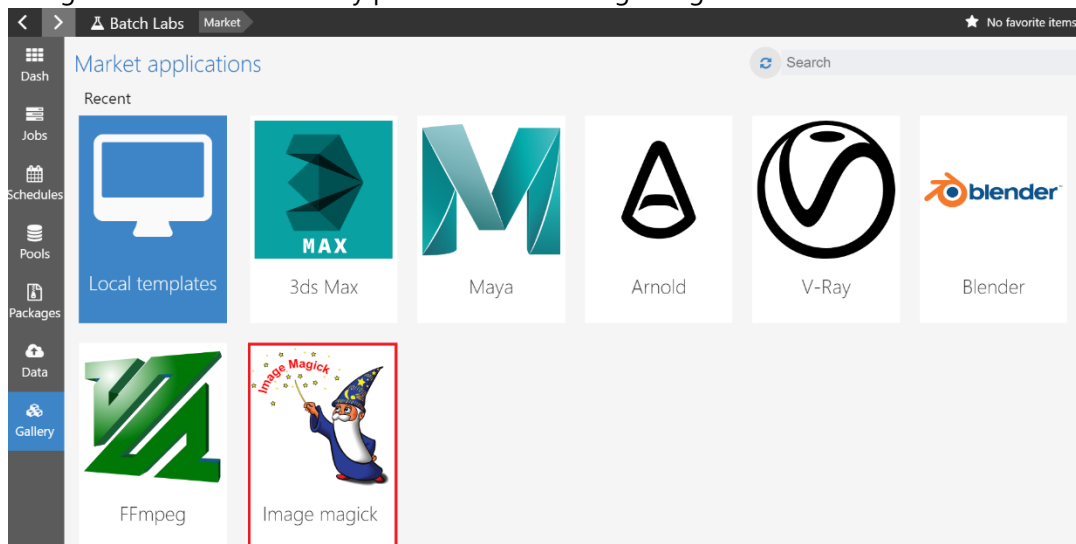
Confirm Cancel

10. Verify that the out file group fgrp-image-magick-output is successfully created:




Step 3: Start Image Magick resize job for provided input files

11. Navigate to Batch Labs Gallery panel and select Image Magick tab:



12. Select action **Resize a collection of images**:

 Action selection for Image magick ([view readme](#))

Resize a collection of images

This job will resize a collection of images in parallel using Image Magick by creating one task per image.

Note: you can click on **view readme** to visit a website with ImageMagick Batch template repository.

13. Select **Run job with auto pool** in (1) Mode selection:

Run resize-images from imagemagick

1 Mode

Run job with auto pool

Run job with existing pool

Create pool for later use

The job will be executed in automatically managed Azure Batch pool – i.e. the pool will be automatically deployed before the job start and deleted after the job is finished.

14. Fill out the remaining configuration parameters:

Run resize-images from imagemagick

1 Mode

Run job with auto pool

Run job with existing pool

Create pool for later use

2 Pool

Pool name
image-magick-pool

The ID of Azure Batch pool

Vm count
4

The number of virtual machines in the Azure Batch pool where the job will run

Vm size
STANDARD_D1

3 Job

Job name
image-magick-resize

The job name

Input filegroup
fgrp-image-magick-input


The file group where the input images are stored. Any non-image files will fail to convert.

Resize
50

Image resize percent

Output filegroup
fgrp-image-magick-output

The file group where outputs will be stored



In (2) **Pool** definition provide a Batch pool name, e.g. **image-magick-pool**, **VM count** equal half the number of input files (in our case **4**) and select **VM size STANDARD_A1** from the list. In (3) **Job** definition provide the job name, e.g. **image-magick-resize**, and select the existing input and output file groups from the drop-down list – **fgrp-image-magick-input** and **fgrp-image-magick-output** respectively. You may leave the **Resize** factor of **50%** or adjust.

15. When configuration is ready, the **Run** button will turn green:



Click on the **Run** button to start the job.

16. Monitor the job until it's finished.

You can observe the job status in the **Jobs** panel:

Id	State	Created	Started	Completed	Exit code
0	active	a few seconds...			
1	active	a few seconds...			
2	active	a few seconds...			
3	active	a few seconds...			
4	active	a few seconds...			
5	active	a few seconds...			
6	active	a few seconds...			
7	active	a few seconds...			


and the status of the pool executing the job on the **Pools** panel:

- idle (0)
- running (0)
- waitingforstarttask (0)
- offline (0)
- preempted (0)
- Transition states
- creating (0)
- starting (4)
- rebooting (0)
- reimaging (0)
- leavingpool (0)
- Error states
- starttaskfailed (0)
- unusable (0)
- unknown (0)

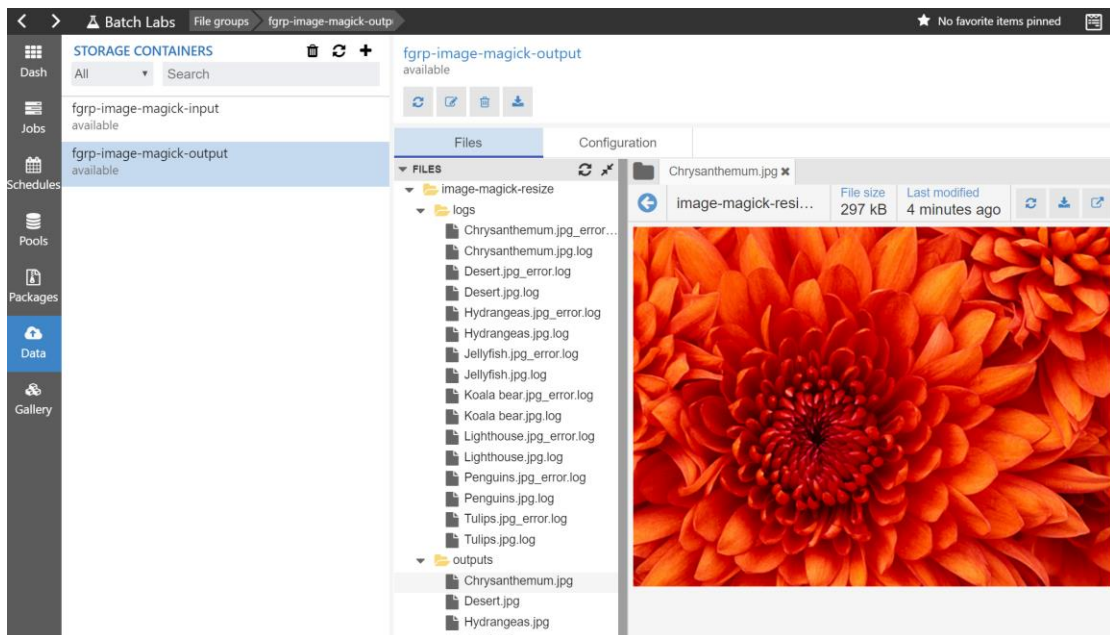
The job should take about 5 minutes to complete.

Step 4: Check result of Image Magick resize job

17. When the job is finished you should see the following status on the Jobs panel (top-right corner):




18. Navigate to Data panel in Batch Labs main window and open fgrp-image-magick-output file group. You should find image-magick-resize job's folder with log and output subfolders:

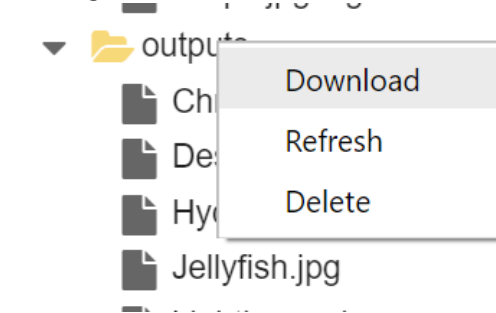


Check that the output jpeg files have reduced sizes as expected.

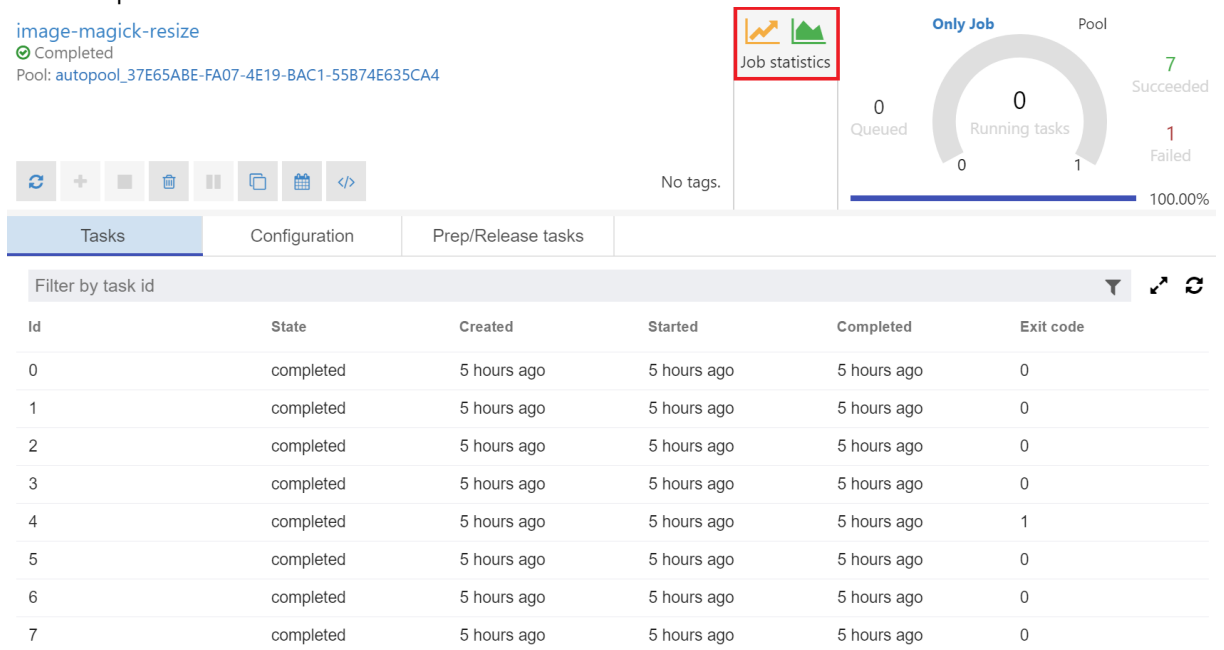
19. You may have observed that one of the tasks failed (for file Koala bear.jpg – due to space character in the name). Check the error message in the log file:




20. You can download the output files to local disk by right-clicking the **output** folder and selecting download action:



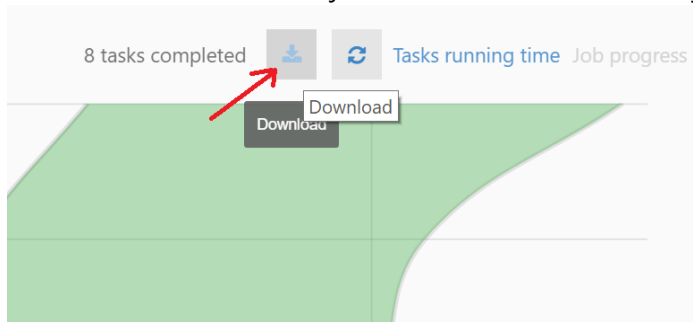
21. You can check job execution statistics by clicking **Job statistics** for *image-magick-resize* job on the **Jobs** panel:



You can view task running times and job progress diagrams:

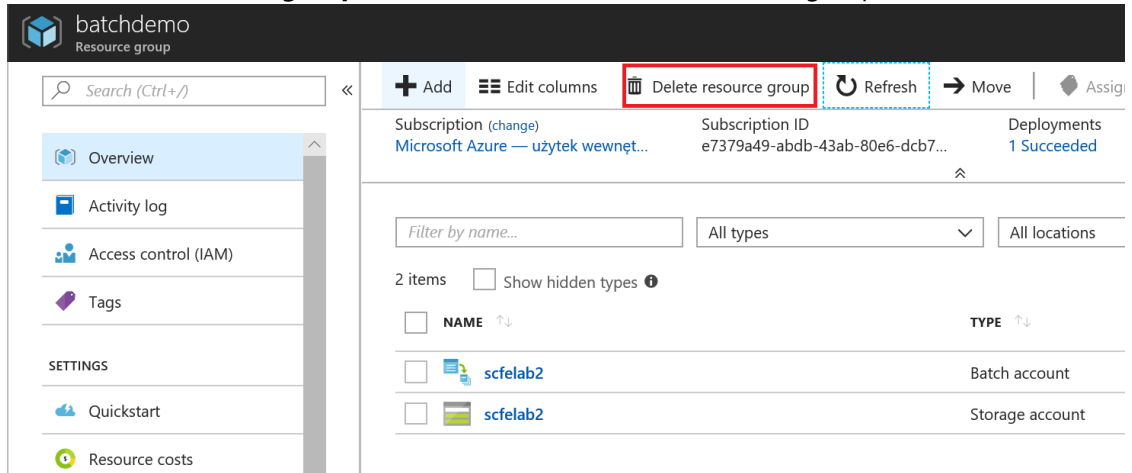


You can also download the job execution data to csv file by clicking download icon:



Step 5: Clean up resources

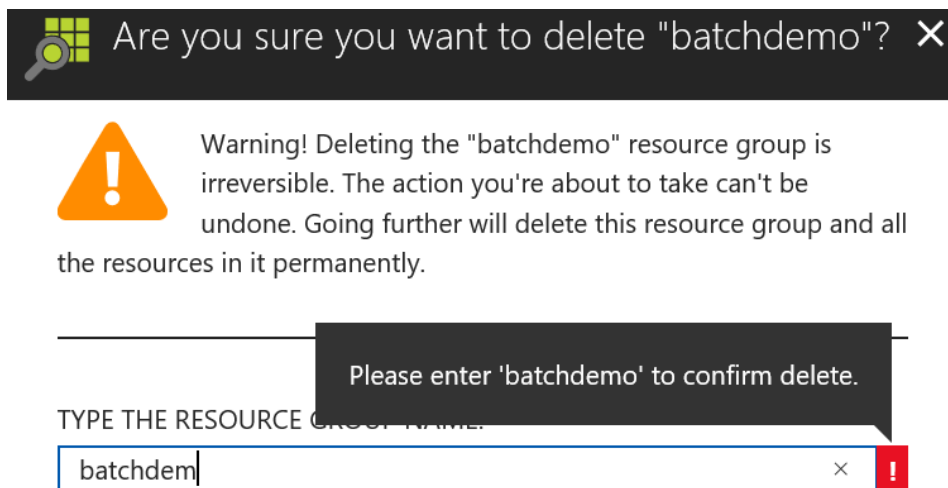
22. As the last step of the exercise clean up the created resources by deleting the resource group. Click **"Delete resource group"** button on the **batchdemo** resource group tab:



The screenshot shows the Azure portal interface for the 'batchdemo' resource group. The left sidebar contains navigation options like Overview, Activity log, Access control (IAM), Tags, and Settings. The main area displays subscription information and a list of resources. The 'Delete resource group' button is highlighted with a red box. Below the button, there are filters for name, type, and location. A table lists two items: 'scfelab2' (Batch account) and 'scfelab2' (Storage account).

NAME	TYPE
scfelab2	Batch account
scfelab2	Storage account

23. Enter the resource group name to confirm resource deletion and click **Delete**:



The screenshot shows a confirmation dialog box with the text: "Are you sure you want to delete 'batchdemo'?". Below the dialog, there is a warning message: "Warning! Deleting the 'batchdemo' resource group is irreversible. The action you're about to take can't be undone. Going further will delete this resource group and all the resources in it permanently." Below the warning, there is a text input field labeled "TYPE THE RESOURCE GROUP NAME:" with the text "batchdem" entered. A red exclamation mark icon is visible next to the input field.